

Machine Learning (CSE 446): Decision Trees

Sham M Kakade

© 2018

University of Washington
`skakade@cs.washington.edu`

Features

Let ϕ be a function that maps from inputs (x) to values.

Features

Let ϕ be a function that maps from inputs (x) to values.

- ▶ If ϕ maps to $\{0, 1\}$, we call it a “binary feature (function).”

Features

Let ϕ be a function that maps from inputs (x) to values.

- ▶ If ϕ maps to $\{0, 1\}$, we call it a “binary feature (function).”
- ▶ If ϕ maps to \mathbb{R} , we call it a “real-valued feature (function).”

Features

Let ϕ be a function that maps from inputs (x) to values.

- ▶ If ϕ maps to $\{0, 1\}$, we call it a “binary feature (function).”
- ▶ If ϕ maps to \mathbb{R} , we call it a “real-valued feature (function).”
- ▶ Feature functions can map to categorical values, ordinal values, integers, and more.

Features

Data derived from <https://archive.ics.uci.edu/ml/datasets/Auto+MPG>

mpg; cylinders; displacement; horsepower; weight; acceleration; year; origin

| | | | | | | | |
|------|---|-------|-------|-------|------|----|---|
| 18.0 | 8 | 307.0 | 130.0 | 3504. | 12.0 | 70 | 1 |
| 15.0 | 8 | 350.0 | 165.0 | 3693. | 11.5 | 70 | 1 |
| 18.0 | 8 | 318.0 | 150.0 | 3436. | 11.0 | 70 | 1 |
| 16.0 | 8 | 304.0 | 150.0 | 3433. | 12.0 | 70 | 1 |
| 17.0 | 8 | 302.0 | 140.0 | 3449. | 10.5 | 70 | 1 |
| 15.0 | 8 | 429.0 | 198.0 | 4341. | 10.0 | 70 | 1 |
| 14.0 | 8 | 454.0 | 220.0 | 4354. | 9.0 | 70 | 1 |
| 14.0 | 8 | 440.0 | 215.0 | 4312. | 8.5 | 70 | 1 |
| 14.0 | 8 | 455.0 | 225.0 | 4425. | 10.0 | 70 | 1 |
| 15.0 | 8 | 390.0 | 190.0 | 3850. | 8.5 | 70 | 1 |
| 15.0 | 8 | 383.0 | 170.0 | 3563. | 10.0 | 70 | 1 |
| 14.0 | 8 | 340.0 | 160.0 | 3609. | 8.0 | 70 | 1 |
| 15.0 | 8 | 400.0 | 150.0 | 3761. | 9.5 | 70 | 1 |
| 14.0 | 8 | 455.0 | 225.0 | 3086. | 10.0 | 70 | 1 |
| 24.0 | 4 | 113.0 | 95.00 | 2372. | 15.0 | 70 | 3 |
| 22.0 | 6 | 198.0 | 95.00 | 2833. | 15.5 | 70 | 1 |
| 18.0 | 6 | 199.0 | 97.00 | 2774. | 15.5 | 70 | 1 |
| 21.0 | 6 | 200.0 | 85.00 | 2587. | 16.0 | 70 | 1 |
| 27.0 | 4 | 97.00 | 88.00 | 2130. | 14.5 | 70 | 3 |
| 26.0 | 4 | 97.00 | 46.00 | 1835. | 20.5 | 70 | 2 |
| 25.0 | 4 | 110.0 | 87.00 | 2672. | 17.5 | 70 | 2 |
| 24.0 | 4 | 107.0 | 90.00 | 2430. | 14.5 | 70 | 2 |

Goal: predict whether mpg is < 23 (“bad” = 0) or above (“good” = 1) given other attributes (other columns).

201 “good” and 197 “bad”;
guessing the most frequent class (good) will get 50.5% accuracy.

Contingency Table

| values of y | values of feature ϕ | | | |
|---------------|--------------------------|-------|----------|-------|
| | v_1 | v_2 | \cdots | v_K |
| | | | | |
| 0 | | | | |
| 1 | | | | |

Decision Stump Example

| y | maker | | |
|-----|---------|--------|------|
| | america | europa | asia |
| 0 | 174 | 14 | 9 |
| 1 | 75 | 56 | 70 |

\downarrow \downarrow \downarrow


0 1 1

Decision Stump Example

| y | maker | | |
|-----|---------|--------|------|
| | america | europa | asia |
| 0 | 174 | 14 | 9 |
| 1 | 75 | 56 | 70 |

\downarrow \downarrow \downarrow

0 1 1




Decision Stump Example

| y | maker | | |
|-----|---------|--------|------|
| | america | europa | asia |
| 0 | 174 | 14 | 9 |
| 1 | 75 | 56 | 70 |


\downarrow \downarrow \downarrow

0 1 1


Errors: $75 + 14 + 9 = 98$ (about 25%)



Decision Stump Example



Decision Stump Example



Errors: $1 + 20 + 1 + 11 + 3 = 36$ (about 9%)


Key Idea: Recursion

A single feature **partitions** the data.

For each partition, we could choose another feature and partition further.


Applying this recursively, we can construct a **decision tree**.

Decision Tree Example




Error reduction compared to the cylinders stump?

Decision Tree Example



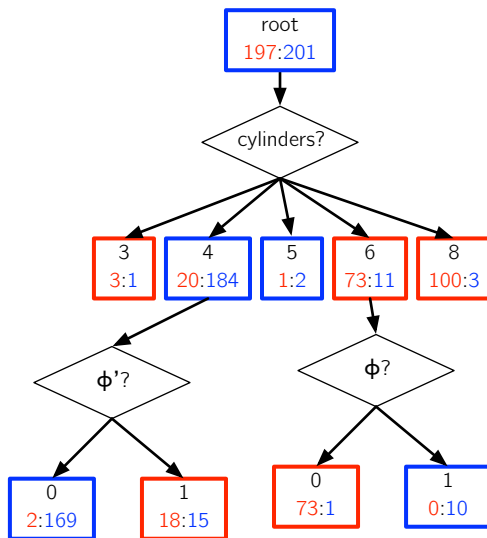
Error reduction compared to the cylinders stump?

Decision Tree Example



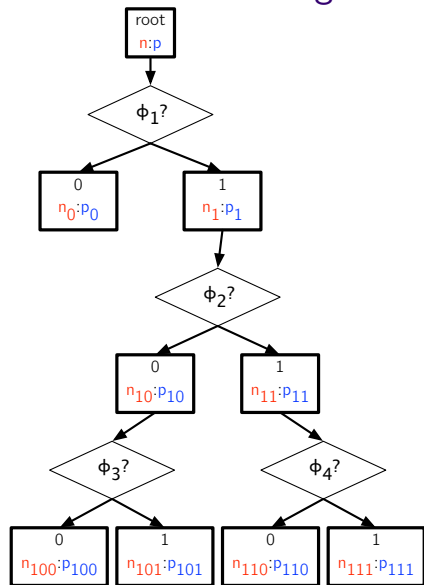
Error reduction compared to the cylinders stump?

Decision Tree Example

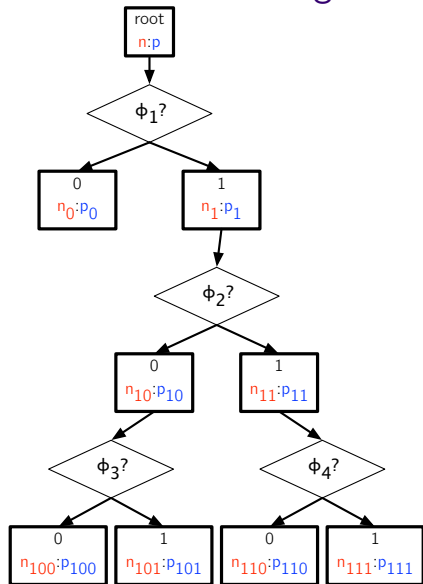


Error reduction compared to the cylinders stump?

Decision Tree: Making a Prediction



Decision Tree: Making a Prediction



Data: decision tree t , input example x

Result: predicted class

if t has the form $\text{LEAF}(y)$ **then**

 return y ;

else

 # $t.\phi$ is the feature associated with t ;


 # $t.\text{child}(v)$ is the subtree for value v ;

 return $\text{DTREETEST}(t.\text{child}(t.\phi(x)), x)$;

end

Algorithm 1: DTREETEST

Decision Tree: Making a Prediction



Equivalent boolean formulas:

$$(\phi_1 = 0) \Rightarrow \llbracket n_0 < p_0 \rrbracket$$

$$(\phi_1 = 1) \wedge (\phi_2 = 0) \wedge (\phi_3 = 0) \Rightarrow \llbracket n_{100} < p_{100} \rrbracket$$

$$(\phi_1 = 1) \wedge (\phi_2 = 0) \wedge (\phi_3 = 1) \Rightarrow \llbracket n_{101} < p_{101} \rrbracket$$

$$(\phi_1 = 1) \wedge (\phi_2 = 1) \wedge (\phi_4 = 0) \Rightarrow \llbracket n_{110} < p_{110} \rrbracket$$

$$(\phi_1 = 1) \wedge (\phi_2 = 1) \wedge (\phi_4 = 1) \Rightarrow \llbracket n_{111} < p_{111} \rrbracket$$

Tangent: How Many Formulas?

Assume we have D binary features.

Tangent: How Many Formulas?

Assume we have D binary features.

Each feature could be set to 0, or set to 1, or excluded (wildcard/don't care).

Tangent: How Many Formulas?

Assume we have D binary features.


Each feature could be set to 0, or set to 1, or excluded (wildcard/don't care).

3^D formulas.

Growing a Decision Tree


root
n:p

Growing a Decision Tree




We chose feature ϕ_1 . Note that $n = n_0 + n_1$ and $p = p_0 + p_1$.

Growing a Decision Tree




We chose not to split the left partition. Why not?

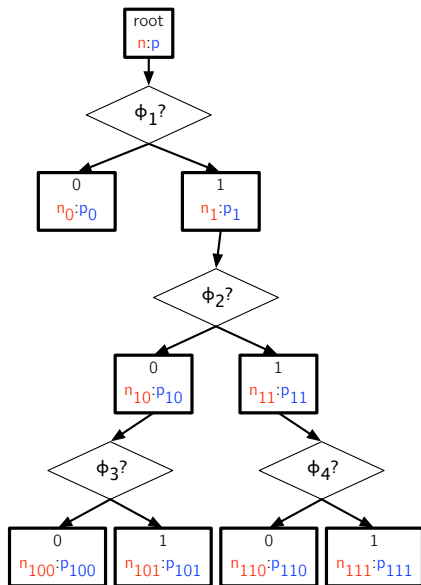
Growing a Decision Tree



Growing a Decision Tree



Growing a Decision Tree



Greedy Building a Decision Tree (Binary Features)

Data: data D , feature set Φ

Result: decision tree

if *all examples in D have the same label y , or Φ is empty and y is the best guess* **then**

 return LEAF(y);

else

for *each feature ϕ in Φ* **do**

 partition D into D_0 and D_1 based on ϕ -values;

 let mistakes(ϕ) = (non-majority answers in D_0) + (non-majority answers in D_1);

end


 let ϕ^* be the feature with the smallest number of mistakes;

 return NODE(ϕ^* , {0 \rightarrow DTREETRAIN(D_0 , $\Phi \setminus \{\phi^*\}$), 1 \rightarrow DTREETRAIN(D_1 , $\Phi \setminus \{\phi^*\}$)});

end

Algorithm 2: DTREETRAIN

Danger: Overfitting



Detecting Overfitting

If you use all of your data to train, you won't be able to draw the red curve on the preceding slide!

Detecting Overfitting

If you use all of your data to train, you won't be able to draw the red curve on the preceding slide!

Solution: hold some out. This data is called **development data**. More terms:

- ▶ Decision tree max depth is an example of a **hyperparameter**
- ▶ “I used my development data to **tune** the max-depth hyperparameter.”

Detecting Overfitting

If you use all of your data to train, you won't be able to draw the red curve on the preceding slide!

Solution: hold some out. This data is called **development data**. More terms:

- ▶ Decision tree max depth is an example of a **hyperparameter**
- ▶ “I used my development data to **tune** the max-depth hyperparameter.”

Better yet, hold out two subsets, one for tuning and one for a true, honest-to-science **test**.

Detecting Overfitting

If you use all of your data to train, you won't be able to draw the red curve on the preceding slide!

Solution: hold some out. This data is called **development data**. More terms:

- ▶ Decision tree max depth is an example of a **hyperparameter**
- ▶ “I used my development data to **tune** the max-depth hyperparameter.”

Better yet, hold out two subsets, one for tuning and one for a true, honest-to-science **test**.

Splitting your data into training/development/test requires careful thinking. Starting point: randomly shuffle examples with an 80%/10%/10% split.